

Approve the Rule, Not the Case

Separating how much oversight a decision needs
from what the decision is

confused.de

ACM Europe School on Responsible AI

The remedy everyone reaches for

A person reviews each consequential decision
before it takes effect.

The remedy everyone reaches for

A person reviews each consequential decision
before it takes effect.

This is called meaningful human control.

It fails in two directions at once

- **At volume it is unaffordable.** So it is skipped, or performed at a rate that makes it ceremonial.
- **Shown a recommendation without its reasoning, a reviewer can only agree.** The check meant to catch errors becomes a signature that launders them.

It fails in two directions at once

- **At volume it is unaffordable.** So it is skipped, or performed at a rate that makes it ceremonial.
- **Shown a recommendation without its reasoning, a reviewer can only agree.** The check meant to catch errors becomes a signature that launders them.

Tuning the threshold manages the cost. It does not touch the cause.

Two questions, one mechanism

How much oversight?

- harm if wrong
- reversibility
- is a right engaged
- how complete the evidence is

A policy over consequences.

What is the decision?

- conditions
- conclusions
- a derivation

Logic.

Two questions, one mechanism

How much oversight?

- harm if wrong
- reversibility
- is a right engaged
- how complete the evidence is

A policy over consequences.

What is the decision?

- conditions
- conclusions
- a derivation

Logic.

Answering both with one mechanism forces the trade-off.

Answering them separately dissolves it.

Half one: the ladder

```
if (harm >= 9 and rights) or (harm >= 8 and reversibility <= 3):  
    human_only          # AI may not decide  
if harm >= 6 or rights or confidence < 70:  
    human_approval      # a person approves first  
if harm <= 3 and reversibility >= 7:  
    ai_with_oversight   # bounded, logged, reversible  
ai_recommends           # propose; do not act
```

- Nine lines, in a module that **imports nothing at all**.
- A module that cannot reach a network cannot be talked into anything.
- It never sees the answer.

The ladder exists twice — so we test that

It runs in Python on the server, and in JavaScript in the page so the sliders respond without a round trip.

- Two copies of the rule allocating human oversight can drift with **nothing failing**.
- The page would show one tier while the server acted on another.

The ladder exists twice — so we test that

It runs in Python on the server, and in JavaScript in the page so the sliders respond without a round trip.

- Two copies of the rule allocating human oversight can drift with **nothing failing**.
- The page would show one tier while the server acted on another.

A test extracts *classify from the shipped page*, runs it under Node, and compares **all 1,000 input combinations**.

Mutating one threshold: twenty disagreements, each one named.

Half two: the logic database

```
will_get_wet(C) ← going_out(C), raining(C),  
                  not has_umbrella(C).
```

- Umbrella recorded **absent** → the rule fires.
- Umbrella **unmentioned** → it does not, and the head is *blocked*, not false.

Half two: the logic database

```
will_get_wet(C) ← going_out(C), raining(C),  
                  not has_umbrella(C).
```

- Umbrella recorded **absent** → the rule fires.
- Umbrella **unmentioned** → it does not, and the head is *blocked*, not false.

Facts come from a model reading free text. Under a closed world, silence becomes evidence against a person one round later.

We shipped that defect twice in the predecessor.

The graph, and what a model may touch

intake → suggest → classify → extract → derive → propose → gate →
reply

Calls a model

- is there a decision here
- estimate the four inputs
- rewrite the case as facts
- draft candidate rules

Does not

- classify — the ladder
- derive — the engine
- gate — is a person needed

The graph, and what a model may touch

intake → suggest → classify → extract → derive → propose → gate →
reply

Calls a model

- is there a decision here
- estimate the four inputs
- rewrite the case as facts
- draft candidate rules

Does not

- classify — the ladder
- derive — the engine
- gate — is a person needed

None of the four produces an answer.

The gate asks one question

Does an approved rule cover this case?

Tier	Approved rule fires	Nothing fires
human_only	proceeds	waits
human_approval	proceeds	waits
ai_with_oversight	proceeds	proceeds, sampled
ai_recommends	proceeds	proceeds, sampled

The gate asks one question

Does an approved rule cover this case?

Tier	Approved rule fires	Nothing fires
human_only	proceeds	waits
human_approval	proceeds	waits
ai_with_oversight	proceeds	proceeds, sampled
ai_recommends	proceeds	proceeds, sampled

The tier sets what is at stake — not whether a person is asked.

Why approved logic proceeds even at the top tier

A rule an expert approved **is their decision**.

They saw the clause. They accepted that any case matching its body gets its conclusion. Applying it executes a call they have made.

Why approved logic proceeds even at the top tier

A rule an expert approved **is their decision**.

They saw the clause. They accepted that any case matching its body gets its conclusion. Applying it executes a call they have made.

- Asking again for every matching case is how oversight becomes ceremony.
- Approval is of a **clause**, whose body states exactly what it claims.
- A proportion is audited anyway — 50% for `human_only` logic, 5% routine, **never zero**.

What broke

- **A greeting was classified.** “Hi” → *human approval required*. The estimator said “no decision described” and scored confidence 0; the ladder read < 70 and demanded a person.
- **Extraction returned prose.** Folded into the estimator’s call to save a request — one prompt, two jobs. Every fact dropped at parse time. Cheaper, and useless.
- **The proposer drafted the opposite rule.** For a disputed debt with no notice, it required `not debt_disputed` and `formal_notice_given`.
- **Parked inquiries contributed nothing.** The most consequential decisions were the ones whose logic was never written down.

The one we got wrong four times

We called the intake gate flaky. Measured, it was **deterministically wrong**:

"Should we reorder toner now?"	6/6 hand over
"Should we reorder envelopes now?"	0/6 hand over

The only difference: *toner* appeared in the prompt's example list.

The one we got wrong four times

We called the intake gate flaky. Measured, it was **deterministically wrong**:

"Should we reorder toner now?"	6/6 hand over
"Should we reorder envelopes now?"	0/6 hand over

The only difference: *toner* appeared in the prompt's example list.

Four rounds of rewording, inconsistent each time → **change the shape, not the words.**

Ask the *narrow* question — was this purely conversational? — and hand over unless yes.

Recognising a greeting is reliable. Recognising a decision is not.

Coverage: the only number that matters

	Run A	Run B
Decided from approved logic	2 / 12	3 / 12
Parked for a person	5 / 12	4 / 12
Coverage	18%	30%
Approved rules	7	5
Distinct predicates coined	19	15

Coverage: the only number that matters

	Run A	Run B
Decided from approved logic	2 / 12	3 / 12
Parked for a person	5 / 12	4 / 12
Coverage	18%	30%
Approved rules	7	5
Distinct predicates coined	19	15

Twelve cases, empty database, every proposed rule approved.

We report the range, not the better number.

Why it is not higher — in the system's own words

The vocabulary the database built for itself:

reorder_toner	reorder_paper
toner_stock_low	paper_stock_low
toner_stock_depleted	paper_stock_depleted

Six terms and two rules for one idea.

Why it is not higher — in the system's own words

The vocabulary the database built for itself:

reorder_toner	reorder_paper
toner_stock_low	paper_stock_low
toner_stock_depleted	paper_stock_depleted

Six terms and two rules for one idea.

- Toner logic cannot decide a paper case → it arrives uncovered → it coins its own near-duplicate.
- Where the entity *stayed out* of the term — `in_arrears`, `debt_disputed` — three cases were decided by the first one's logic.

`reorder(C, X) ← stock_low(C, X), not depleted(C, X).`

The engine is first-order and can already express this. Nothing proposes it.

What we have not measured

- **No expert has used this.** Every approval in those runs was granted by a script.
- **One model.** All figures are Haiku 4.5. Whether a stronger model coins a reusable vocabulary is the bottleneck, untested and cheap to test.
- **Two runs is not a measurement.** 18 vs 30 per cent on identical input.
- **Top-tier conduct is a prompt, not a mechanism.** It held under direct pressure in every trial — three trials is an anecdote.

Per-decision oversight degrades into ceremony.

Rule-level oversight amortises —
but only if rules generalise, and ours do not yet.

A fifth of cases, decided the same way every time,
with a derivation attached and a named person behind the logic.

The path from a fifth to a majority runs through one mechanism
we can describe, have not built, and can see the shape of
in our own system's vocabulary.