

Table 3 Topic: Human-Centred, Agentic, Decision-Support AI

One-sentence pitch: *We want to move beyond simply putting a “human in the loop” and determine what meaningful human control actually looks like when an increasing number of autonomous AI agents participate in consequential decisions.*

From AI agent to accountable decision partner, how should agentic AI support consequential human decisions while preserving meaningful human agency and accountability and responsible oversight?

Objectives — outputs and outcomes.

Outputs: (1) a simple **AI–Human Decision Rights Matrix** classifying tasks as *AI recommend / AI act with oversight / human approval required / human-only*,

(2) a prototype decision-support/agent workflow incorporating escalation, explanation and human override; and

(3) a short set of design principles for responsible agentic AI. **Outcomes:** better-calibrated human trust; reduced automation bias and over-reliance; clearer accountability; demonstrably meaningful rather than tokenistic human oversight; and a reusable approach applicable across high-consequence settings.

Layer	Description		
1	Problem Domain		
2	Generalisable Model of Human to AI control		
3	Full AI Decision Control	AI-Human Mixed Decision Control	Full Human Decision Control

Define metrics measuring success.

Test the prototype/model against a baseline using **decision quality** (accuracy/error rate), **human agency** (% of inappropriate AI recommendations correctly challenged), **trust calibration** (whether confidence in AI corresponds to its actual reliability), **explainability/comprehension**, **fairness** across relevant groups, **accountability** (can responsibility for each decision be identified?), and **efficiency** (time/cognitive workload). A particularly strong metric would be the **appropriate reliance rate**: *do humans accept the AI when it is right and challenge it when it is wrong?*

For **Group 3: Human-Centred, Agentic & Decision-Support AI**, I would frame the project around a problem that is both academically credible for an ACM Responsible AI School and practically demonstrable.

- **Propose a topic — “From AI Agent to Accountable Decision Partner”**
How should agentic AI support consequential human decisions while preserving meaningful human agency, accountability and appropriate oversight? The project would develop a human-centred model for deciding **what AI may do autonomously, what requires human approval, and what should remain**

exclusively human. A practical scenario—such as procurement, recruitment, healthcare or another public-service decision—could be used to test the model.

- **Draft a timeline**

Week 1: define the decision context, stakeholders, harms and literature; **Week 2:** map human/AI decision rights and responsible-AI requirements; **Weeks 3–4:** design the matrix and prototype workflow; **Weeks 5–6:** construct scenarios, including deliberately erroneous, biased or uncertain AI recommendations; **Weeks 7–8:** conduct user evaluation and collect quantitative/qualitative evidence; **Week 9:** refine the model from findings; **Week 10:** finalise the framework, prototype/demo and recommendations.